Container Network - Linux Networking Basics

NETWORK NAMESPACE INSPECTING

Linux Network Stack
network-stack

In Linux networking, within the Linux network stack, routes define traffic paths, iptables configures packet filtering, lo is a local loopback interface for testing, and eth0 is the primary Ethernet interface for external connections. Let's inspect the network stack, in short.

Network interfaces allow us to establish communication between a network and a device.

ip link list
lo is the loopback interface, allowing local network communication within a device without external network involvement. Verify the loopback interface is up

ifconfig lo
A route in networking specifies the path for network traffic from source to destination. View the routing table:

ip route show
iptables is a user-space utility for configuring packet filter rules in the Linux kernel's Netfilter framework. View iptables rules:

iptables -L
Create Custom Network Namespace
Let's create a custom namespace using ip netns add utiliy.

sudo ip netns add poridhi
sudo ip netns list
Now, entering a network namespace in Linux:

ip netns exec

It is part of the iproute2 package and is often used for managing network namespaces.

sudo ip netns exec poridhi bash
nsenter

The nsenter utility is commonly used to enter into namespaces in Linux, including network namespaces.

sudo nsenter --net=/var/run/netns/poridhi bash
Now, check the network interfaces inside the new ns.

ip link show
Let's check for iptable rules for custom ns.

ifconfig lo

explain these for a simple child in a very easiest manner in details with basics also correct any error

Sure! Let's break everything down in the easiest way possible. Imagine your computer is like a big city, and different parts of the city (apps and services) need to talk to each other. The **Linux Network Stack** is like the system of roads, traffic lights, and highways that let information move from one place to another.

Understanding the Basics of Networking in Linux

Think of a **network interface** like a door to your house. It lets information (data) come in and go out.

1. Listing Network Interfaces

To see all the doors (network interfaces) in your computer, use:

```bash
ip link list
```

This will show you all the available network interfaces.

- **lo (Loopback Interface)**: This is a special door inside your own house. It lets programs on your computer talk to each other **without using the internet**.
- **eth0 (Ethernet Interface)**: This is like the main gate of your house. It lets you connect to the outside world (like the internet).

To check if the **loopback interface** (lo) is working, use:

```bash
ifconfig lo
```

If it's up, it means your computer can communicate with itself.

---

2. Viewing the Routes (Roads for Data)

In a city, roads guide where cars go. In a computer network, **routes** guide where data should go.

To see the list of roads (routes), use:

```bash
ip route show
```

This tells the computer how to send and receive messages.

---

3. Understanding iptables (Security Guards for Traffic)

Imagine security guards at checkpoints controlling what enters and leaves the city. **iptables** is a tool that acts like these security guards, deciding which data packets can pass through.

To check the security rules, use:

```
bash

iptables -L
```

This will list all the rules currently in place.

Creating a Separate "City" (Network Namespace)

A **network namespace** is like creating a new city inside your computer. It will have its own separate roads, doors, and security guards.

To create a new **network namespace**, use:

```
bash

sudo ip netns add poridhi
```

Now, to check if it was created, use:

```
bash

sudo ip netns list
```

It will show a list of all network namespaces, including "poridhi".

Entering the New Namespace (Visiting the New City)

Once we create a new network namespace, we need to **enter** it to check things inside.

Method 1: Using `**ip netns exec**`

```
bash

sudo ip netns exec poridhi bash
```

This command lets us "step inside" the new city (namespace).

Method 2: Using `**nsenter**`

```
bash

sudo nsenter --net=/var/run/netns/poridhi bash
```

This is another way to enter the new city.

Checking the Network Inside the New Namespace

Now that we're inside the new network namespace, let's check what network interfaces exist.

```
bash

ip link show
```

Since it's a new network space, it may not have `**eth0**` yet.

To check if the loopback interface is working inside the new namespace:

```
bash

ifconfig lo
```

If it's not up, you may need to turn it on:

```
bash

ip link set lo up
```

Summary

1. `**ip link list**` → Shows all network interfaces (doors).
2. `**ifconfig lo**` → Checks if the computer can talk to itself.
3. `**ip route show**` → Shows data paths (roads).
4. `**iptables -L**` → Shows security rules.
5. `**ip netns add poridhi**` → Creates a separate network space.
6. `**sudo ip netns list**` → Lists all network spaces.
7. `**sudo ip netns exec poridhi bash**` → Enters the new network space.
8. `**ip link show**` → Checks network interfaces inside the new space.
9. `**ifconfig lo**` → Checks loopback inside the namespace.

This is like setting up a small city inside your computer, separate from the main one. 🚀

explain these terminals:
brd 00:00:00:00:00:00
2: bond0: <BROADCAST,MULTICAST,MASTE
R> mtu 1500 qdisc noqueue state DOWN
mode DEFAULT group default qlen 100
0

```
    link/ether d2:25:4b:d2:1e:34 brd
 ff:ff:ff:ff:ff:ff
3: dummy0: <BROADCAST,NOARP> mtu 150
0 qdisc noop state DOWN mode DEFAULT
 group default qlen 1000
    link/ether 6e:6f:ec:6d:af:a8 brd
 ff:ff:ff:ff:ff:ff
4: eth0: <BROADCAST,MULTICAST,UP,LOW
ER_UP> mtu 1500 qdisc pfifo_fast sta
te UP mode DEFAULT group default qle
n 1000
    link/ether b6:b1:28:ec:3e:1f brd
 ff:ff:ff:ff:ff:ff
5: docker0: <NO-CARRIER,BROADCAST,MU
LTICAST,UP> mtu 1500 qdisc noqueue s
tate DOWN mode DEFAULT group default

    link/ether 02:42:b0:d6:bc:97 brd
 ff:ff:ff:ff:ff:ff
root@50cd87437a310287:~# ifconfig lo
lo: flags=73<UP,LOOPBACK,RUNNING>  m
tu 65536
        inet 127.0.0.1  netmask 255.
0.0.0
        inet6 ::1  prefixlen 128  sc
opeid 0x10<host>
        loop  txqueuelen 1000  (Loca
l Loopback)
        RX packets 0  bytes 0 (0.0 B
)
        RX errors 0  dropped 0  over
runs 0  frame 0
        TX packets 0  bytes 0 (0.0 B
)
        TX errors 0  dropped 0 overr
uns 0  carrier 0  collisions 0
root@50cd87437a310287:~# ip route show
default via 10.62.0.1 dev eth0
10.62.0.0/16 dev eth0 proto kernel scope link src 10.62.18.251
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1 linkdown
root@50cd87437a310287:~# iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source               destination
Chain FORWARD (policy DROP)
target     prot opt source               destination
DOCKER-USER  all  --  anywhere             anywhere
DOCKER-ISOLATION-STAGE-1  all  --  anywhere             anywhere
ACCEPT     all  --  anywhere             anywhere             ctstate RELATED,
ESTABLISHED
DOCKER     all  --  anywhere             anywhere
ACCEPT     all  --  anywhere             anywhere
```

```
ACCEPT     all  --  anywhere             anywhere
Chain OUTPUT (policy ACCEPT)
target     prot opt source               destination
Chain DOCKER (1 references)
target     prot opt source               destination
Chain DOCKER-ISOLATION-STAGE-1 (1 references)
target     prot opt source               destination
DOCKER-ISOLATION-STAGE-2  all  --  anywhere             anywhere
RETURN     all  --  anywhere             anywhere
Chain DOCKER-ISOLATION-STAGE-2 (1 references)
target     prot opt source               destination
DROP       all  --  anywhere             anywhere
RETURN     all  --  anywhere             anywhere
Chain DOCKER-USER (1 references)
target     prot opt source               destination
RETURN     all  --  anywhere             anywhere
root@50cd87437a310287:~#
```

Let's break down the terminal outputs into simple explanations.

---

1. Understanding the Network Interfaces

When you run `**`ip link list`**`, it shows the network interfaces on your system. Here's what each one means:

brd 00:00:00:00:00:00

- `**`brd`**` stands for **broadcast address**.
- `**`00:00:00:00:00:00`**` means it's not set (default).
- **Broadcasting** is when a network message is sent to all devices in a network.

2: bond0

```
bash

2: bond0: <BROADCAST,MULTICAST,MASTER> mtu 1500 qdisc noqueue state DOWN mode DEFAULT group default
    link/ether d2:25:4b:d2:1e:34 brd ff:ff:ff:ff:ff:ff
```

- **bond0**: This is a **bonded interface**, meaning it combines multiple network connections for reliability or speed.
- **BROADCAST, MULTICAST, MASTER**:
  - **BROADCAST**: Can send packets to all devices in the network.
  - **MULTICAST**: Can send packets to a group of devices.
  - **MASTER**: This interface manages bonded network connections.
- **MTU 1500**: Maximum packet size is 1500 bytes.
- **State DOWN**: It is currently inactive.

- **MAC Address**: `d2:25:4b:d2:1e:34` (unique hardware identifier).

3: dummy0

```
bash

3: dummy0: <BROADCAST,NOARP> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/ether 6e:6f:ec:6d:af:a8 brd ff:ff:ff:ff:ff:ff
```

- **dummy0**: A **dummy network interface**, often used for testing or software purposes.
- **NOARP**: This interface does **not use ARP** (Address Resolution Protocol).
- **State DOWN**: It is currently inactive.

4: eth0

```
bash

4: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT group de
    link/ether b6:b1:28:ec:3e:1f brd ff:ff:ff:ff:ff:ff
```

- **eth0**: This is the main **Ethernet network interface** (wired internet).
- **UP, LOWER_UP**: The network interface is active and connected.
- **pfifo_fast**: A queuing discipline used to manage network traffic.

5: docker0

```
bash

5: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT grou
    link/ether 02:42:b0:d6:bc:97 brd ff:ff:ff:ff:ff:ff
```

- **docker0**: This is a virtual network bridge created by **Docker** to connect containers.
- **NO-CARRIER**: No physical network is connected to it.
- **State DOWN**: It is not currently in use.

---

2. Loopback Interface (`lo`)

```
bash

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
        inet 127.0.0.1  netmask 255.0.0.0
        inet6 ::1  prefixlen 128  scopeid 0x10<host>
```

- **lo (Loopback Interface)**: A virtual interface used for internal communication **within the same computer**.
- **IP Address 127.0.0.1**: Used to send data to itself (localhost).
- **MTU 65536**: This allows very large packets (much bigger than normal interfaces).

3. Network Routing Table (`ip route show`)

```
bash

default via 10.62.0.1 dev eth0
10.62.0.0/16 dev eth0 proto kernel scope link src 10.62.18.251
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1 linkdown
```

- **default via 10.62.0.1 dev eth0**: All outgoing internet traffic is sent through `eth0` (wired internet).
- **10.62.0.0/16 dev eth0**: The computer is part of a local network (subnet `10.62.0.x`).
- **172.17.0.0/16 dev docker0**: This is Docker's internal network.
  - `linkdown`: Docker network is not active.

---

4. Firewall Rules (`iptables -L`)

This is a list of firewall rules that control network traffic.

Chain INPUT (Incoming Traffic)

```
bash

Chain INPUT (policy ACCEPT)
target     prot opt source               destination
```

- **Policy ACCEPT**: No restrictions for incoming traffic.

Chain FORWARD (Traffic Routing)

```
bash

Chain FORWARD (policy DROP)
```

- **Policy DROP**: By default, the system **blocks** forwarded traffic (useful in routers).
- Other rules:

```
bash

DOCKER-USER  all  --  anywhere             anywhere
DOCKER-ISOLATION-STAGE-1  all  --  anywhere             anywhere
ACCEPT     all  --  anywhere             anywhere             ctstate RELATED, ESTABLISHED
```

- **DOCKER Rules**: Docker has its own firewall rules to isolate container traffic.
- **ACCEPT all RELATED, ESTABLISHED**: Allow existing connections.

Chain OUTPUT (Outgoing Traffic)

```
bash

Chain OUTPUT (policy ACCEPT)
```

- The system allows all outgoing traffic by default.

Docker-Specific Chains

```
bash

Chain DOCKER-ISOLATION-STAGE-1
Chain DOCKER-ISOLATION-STAGE-2
```

- These rules **isolate** Docker containers so they don't interfere with each other.

Summary

1. **`ip link list`** → Lists all network interfaces.
2. **Loopback (`lo`)** → Used for internal communication (**`127.0.0.1`**).
3. **Bonded Interface (`bond0`)** → Merges multiple connections.
4. **Dummy Interface (`dummy0`)** → Fake network interface (for testing).
5. **Ethernet Interface (`eth0`)** → Main internet connection.
6. **Docker Bridge (`docker0`)** → Connects containers.
7. **`ip route show`** → Shows network routes.
8. **`iptables -L`** → Shows firewall rules.

This setup is common on Linux machines running Docker or virtualized environments! 🚀

explain these terminal and how to execute command inside a custom namespaces:
root@50cd87437a310287:~# ip route show
root@50cd87437a310287:~# sudo ip netns add poridhi
Cannot create namespace file "/run/netns/poridhi": File exists
root@50cd87437a310287:~# sudo ip netns list
poridhi
i bash0cd87437a310287:~# sudo ip netns exec poridhi
root@50cd87437a310287:~# ip link show
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN mode DEFAULT group default q
len 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
i bash0cd87437a310287:~# sudo ip netns exec poridhi
i bash0cd87437a310287:~# sudo ip netns exec poridhi
root@50cd87437a310287:~# ifconfig lo
lo: flags=8<LOOPBACK> mtu 65536
        loop txqueuelen 1000 (Local Loopback)

```
        RX packets 0  bytes 0 (0.0 B)
        RX errors 0  dropped 0  overruns 0  frame 0
        TX packets 0  bytes 0 (0.0 B)
        TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0
root@50cd87437a310287:~#
```

Let's go through your terminal commands step by step and explain everything in simple terms.

---

1. Checking Routing Table

```
bash

root@50cd87437a310287:~# ip route show
```

- This command shows the **routing table**, which tells the system where to send network traffic.
- It helps understand how data flows inside and outside your machine.

---

2. Creating a Network Namespace

```
bash

root@50cd87437a310287:~# sudo ip netns add poridhi
Cannot create namespace file "/run/netns/poridhi": File exists
```

- **`sudo ip netns add poridhi`** → This command tries to create a network namespace named **`poridhi`**.
- The error **"File exists"** means the namespace **`poridhi`** **already exists**, so you don't need to create it again.

---

3. Listing Available Network Namespaces

```
bash

root@50cd87437a310287:~# sudo ip netns list
poridhi
```

- This confirms that **`poridhi`** already exists.
- It shows a list of all created **network namespaces**.

---

4. Executing a Command Inside the Namespace

```
bash

root@50cd87437a310287:~# sudo ip netns exec poridhi
```

- **`ip netns exec poridhi`** → This command runs a command **inside the namespace `poridhi`**.

- However, in your output, there's no command after `exec poridhi`, so nothing happens.
- To enter the namespace properly, you should use:

```
bash

sudo ip netns exec poridhi bash
```

This will start a new **Bash shell** inside the `poridhi` namespace.

---

5. Checking Network Interfaces Inside the Namespace

```
bash

root@50cd87437a310287:~# ip link show
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

- `ip link show` lists all network interfaces inside the current namespace.
- Since we haven't created any custom interfaces inside `poridhi`, the only interface present is:
  - **lo (Loopback Interface)**: Used for internal communication within the namespace.
  - **State DOWN**: It is currently disabled.
- To enable the loopback interface inside `poridhi`, use:

```
bash

sudo ip netns exec poridhi ip link set lo up
```

---

6. Checking the Loopback Interface (`lo`)

```
bash

root@50cd87437a310287:~# ifconfig lo
lo: flags=8<LOOPBACK>  mtu 65536
        loop  txqueuelen 1000  (Local Loopback)
        RX packets 0  bytes 0 (0.0 B)
        TX packets 0  bytes 0 (0.0 B)
```

- `ifconfig lo` shows the status of the `lo` (loopback) interface.
- **Flags** `<LOOPBACK>` → This confirms it is a loopback interface.
- **MTU 65536** → Allows large packets inside the namespace.
- **RX and TX packets 0** → No data has been sent or received yet.

How to Properly Execute Commands Inside a Namespace

Entering the Namespace with a Bash Shell

```
bash

sudo ip netns exec poridhi bash
```

- This will start a new Bash session **inside** the namespace.

Running Commands Directly Inside the Namespace

If you just want to run a single command inside `poridhi`, use:

```
bash

sudo ip netns exec poridhi <command>
```

Examples:

- **Check network interfaces inside `poridhi`:**

  ```
  bash

  sudo ip netns exec poridhi ip link show
  ```

- **Enable the loopback interface inside `poridhi`:**

  ```
  bash

  sudo ip netns exec poridhi ip link set lo up
  ```

- **Check the IP address inside `poridhi`:**

  ```
  bash

  sudo ip netns exec poridhi ip addr show
  ```

- **Ping another IP from inside `poridhi`:**

  ```
  bash

  sudo ip netns exec poridhi ping 8.8.8.8
  ```

Summary

1. **Check network routes:** `ip route show`
2. **Create a namespace:** `sudo ip netns add poridhi`
3. **List namespaces:** `sudo ip netns list`
4. **Enter namespace shell:** `sudo ip netns exec poridhi bash`
5. **Run commands inside namespace:**
   - `sudo ip netns exec poridhi ip link show` → List interfaces

- `sudo ip netns exec poridhi ip link set lo up` → Enable loopback
- `sudo ip netns exec poridhi ping 8.8.8.8` → Ping from namespace

Let me know if you need more clarification! 🚀